

# Session 3: YARN

## Assignment

**Note:** The VM available with me is having hadoop 2.6.5 and not the 2.9.0 so executing hadoop-mapreduce-examples 2.6.5 jar instead of 2.9.0 jars.

### Task 1:

- Execute WordMean Job

**Command** - yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /input /output1

The **first highlighted** part indicates the path of the jar file to be executed.

The **second highlighted** part indicates the driver (main) class of the jar file. If main class is not defined while creating a jar file we have to specify the class name in command else we can omit the same.

The **third highlighted** part indicates input folder to be read.

The **forth highlighted** part indicates output folder to be written with output of job.

Below screen shot after executing the command

```
[acadgild@localhost ~]$ yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /input /output1
18/03/24 12:21:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/24 12:21:46 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/03/24 12:21:49 INFO input.FileInputFormat: Total input paths to process : 1
18/03/24 12:21:50 INFO mapreduce.JobSubmitter: number of splits:1
18/03/24 12:21:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1521859699712_0002
18/03/24 12:21:51 INFO impl.YarnClientImpl: Submitted application application_1521859699712_0002
18/03/24 12:21:51 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1521859699712_0002/
18/03/24 12:21:51 INFO mapreduce.Job: Running job: job_1521859699712_0002
18/03/24 12:22:08 INFO mapreduce.Job: Job job_1521859699712_0002 running in uber mode : false
18/03/24 12:22:08 INFO mapreduce.Job: map 0% reduce 0%
18/03/24 12:22:18 INFO mapreduce.Job: map 100% reduce 0%
18/03/24 12:22:31 INFO mapreduce.Job: map 100% reduce 100%
18/03/24 12:22:32 INFO mapreduce.Job: Job job_1521859699712_0002 completed successfully
18/03/24 12:22:32 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215671
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=196
    HDFS: Number of bytes written=19
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=8357
    Total time spent by all reduces in occupied slots (ms)=9191
    Total time spent by all map tasks (ms)=8357
    Total time spent by all reduce tasks (ms)=9191
    Total vcore-milliseonds taken by all map tasks=8357
    Total vcore-milliseonds taken by all reduce tasks=9191
    Total megabyte-milliseonds taken by all map tasks=8557568
    Total megabyte-milliseonds taken by all reduce tasks=9411584
  Map-Reduce Framework
    Map input records=1
    Map output records=34
    Map output bytes=493
    Map output materialized bytes=39
    Input split bytes=101
    Combine input records=34
    Combine output records=2
    Reduce input groups=2
    Reduce shuffle bytes=39
    Reduce input records=2
```

```

        Spilled Records=4
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=230
        CPU time spent (ms)=2400
        Physical memory (bytes) snapshot=295473152
        Virtual memory (bytes) snapshot=4118224896
        Total committed heap usage (bytes)=170004480

    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0

    File Input Format Counters
        Bytes Read=95
    File Output Format Counters
        Bytes Written=19
The mean is: 4.588235294117647
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

Folder got created as output1

```

[acadgild@localhost ~]$ hadoop fs -ls /
18/03/24 12:29:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 9 items
drwxr-xr-x - acadgild supergroup 0 2018-03-04 22:22 /Prachi
drwxr-xr-x - acadgild supergroup 0 2018-03-24 09:38 /hadoopdata
drwxr-xr-x - acadgild supergroup 0 2018-03-04 19:15 /input
drwxr-xr-x - acadgild supergroup 0 2018-03-05 09:06 /output
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:22 /output1
drwxr-xr-x - acadgild supergroup 0 2018-03-17 09:50 /sqoopout
drwxr-xr-x - acadgild supergroup 0 2018-02-02 12:49 /sqoopout111
drwxrwx-- - acadgild supergroup 0 2018-03-05 08:40 /tmp
drwxr-xr-x - acadgild supergroup 0 2018-03-18 07:10 /user
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

[acadgild@localhost ~]$ hadoop fs -ls /output1
18/03/24 12:30:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-03-24 12:22 /output1/ SUCCESS
-rw-r--r-- 1 acadgild supergroup 19 2018-03-24 12:22 /output1/part-r-00000
[acadgild@localhost ~]$
```

Output of word mean

```

[acadgild@localhost ~]$ hadoop fs -cat /output1/part-r-00000
18/03/24 12:30:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count 17
length 78
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

Task 2:

Execute WordMedian

Command - yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /input/output2

Note: The color used for highlights indicates same meaning as explained above.

```
[acadgild@localhost ~]$ yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /input /output2
18/03/24 12:31:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/24 12:32:00 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/03/24 12:32:02 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner
edy this.
18/03/24 12:32:03 INFO input.FileInputFormat: Total input paths to process : 1
18/03/24 12:32:04 INFO mapreduce.JobSubmitter: number of splits:1
18/03/24 12:32:04 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1521859699712_0003
18/03/24 12:32:05 INFO impl.YarnClientImpl: Submitted application application_1521859699712_0003
18/03/24 12:32:05 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1521859699712_0003/
18/03/24 12:32:05 INFO mapreduce.Job: Running job: job_1521859699712_0003
18/03/24 12:32:21 INFO mapreduce.Job: Job job_1521859699712_0003 running in uber mode : false
18/03/24 12:32:21 INFO mapreduce.Job: map 0% reduce 0%
18/03/24 12:32:31 INFO mapreduce.Job: map 100% reduce 0%
18/03/24 12:32:44 INFO mapreduce.Job: map 100% reduce 100%
18/03/24 12:32:44 INFO mapreduce.Job: Job job_1521859699712_0003 completed successfully
18/03/24 12:32:45 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=66
  FILE: Number of bytes written=215461
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=196
  HDFS: Number of bytes written=24
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=8370
  Total time spent by all reduces in occupied slots (ms)=9433
  Total time spent by all map tasks (ms)=8370
  Total time spent by all reduce tasks (ms)=9433
  Total vcore-milliseconds taken by all map tasks=8370
  Total vcore-milliseconds taken by all reduce tasks=9433
  Total megabyte-milliseconds taken by all map tasks=8570880
  Total megabyte-milliseconds taken by all reduce tasks=9659392

Map-Reduce Framework
  Map input records=1
  Map output records=17
  Map output bytes=136
  Map output materialized bytes=66
  Input split bytes=101
  Combine input records=17
  Combine output records=6
```

```
Reduce input records=0
Reduce output records=6
Spilled Records=12
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=249
CPU time spent (ms)=2450
Physical memory (bytes) snapshot=295727104
Virtual memory (bytes) snapshot=4118224896
Total committed heap usage (bytes)=170004480

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=95

File Output Format Counters
  Bytes Written=24

The median is: 4
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

Output2 folder is created

```
[acadgild@localhost ~]$ hadoop fs -ls /
18/03/24 12:34:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 10 items
drwxr-xr-x - acadgild supergroup 0 2018-03-04 22:22 /Prachi
drwxr-xr-x - acadgild supergroup 0 2018-03-24 09:38 /hadoopdata
drwxr-xr-x - acadgild supergroup 0 2018-03-04 19:15 /input
drwxr-xr-x - acadgild supergroup 0 2018-03-05 09:06 /output
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:22 /output1
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:32 /output2
drwxr-xr-x - acadgild supergroup 0 2018-03-17 09:59 /sqoopout
drwxr-xr-x - acadgild supergroup 0 2018-02-02 12:49 /sqoopout111
drwxrwx-- - acadgild supergroup 0 2018-03-05 08:40 /tmp
drwxr-xr-x - acadgild supergroup 0 2018-03-18 07:10 /user
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

```

[acadgild@localhost ~]$ hadoop fs -ls /output2
18/03/24 12:35:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup          0 2018-03-24 12:32 /output2/_SUCCESS
-rw-r--r-- 1 acadgild supergroup        24 2018-03-24 12:32 /output2/part-r-00000
[acadgild@localhost ~]$ hadoop fs -cat /output2/part-r-00000
18/03/24 12:35:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2      4
3      1
4      4
5      3
7      4
8      1
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

## Task 3:

### Execute WordStandardDeviation

Command - yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /input /output3

Note: The color used for highlights indicates same meaning as explained above.

```

[acadgild@localhost ~]$ yarn jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /input /output3
18/03/24 12:36:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/24 12:36:48 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/03/24 12:36:51 INFO input.FileInputFormat: Total input paths to process : 1
18/03/24 12:36:52 INFO mapreduce.JobSubmitter: number of splits:1
18/03/24 12:36:52 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1521859699712_0004
18/03/24 12:36:53 INFO impl.YarnClientImpl: Submitted application application_1521859699712_0004
18/03/24 12:36:53 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1521859699712_0004/
18/03/24 12:36:53 INFO mapreduce.Job: Running job: job_1521859699712_0004
18/03/24 12:37:08 INFO mapreduce.Job: Job job_1521859699712_0004 running in uber mode : false
18/03/24 12:37:08 INFO mapreduce.Job:  map 0% reduce 0%
18/03/24 12:37:19 INFO mapreduce.Job:  map 100% reduce 0%
18/03/24 12:37:32 INFO mapreduce.Job:  map 100% reduce 100%
18/03/24 12:37:33 INFO mapreduce.Job: Job job_1521859699712_0004 completed successfully
18/03/24 12:37:33 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=56
    FILE: Number of bytes written=215869
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=196
    HDFS: Number of bytes written=30
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=8547
    Total time spent by all reduces in occupied slots (ms)=9346
    Total time spent by all map tasks (ms)=8547
    Total time spent by all reduce tasks (ms)=9346
    Total vcore-milliseconds taken by all map tasks=8547
    Total vcore-milliseconds taken by all reduce tasks=9346
    Total megabyte-milliseconds taken by all map tasks=8752128
    Total megabyte-milliseconds taken by all reduce tasks=9570304
  Map-Reduce Framework
    Map input records=1
    Map output records=51
    Map output bytes=748
    Map output materialized bytes=56
    Input split bytes=101
    Combine input records=51
    Combine output records=3
    Reduce input groups=3
    Reduce shuffle bytes=56
    Reduce input records=3

```

```

        Spilled Records=6
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=222
        CPU time spent (ms)=2510
        Physical memory (bytes) snapshot=295038976
        Virtual memory (bytes) snapshot=4118224896
        Total committed heap usage (bytes)=170004480
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=95
    File Output Format Counters
        Bytes Written=30
The standard deviation is: 1.9721240722635547
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ █

```

Output3 folder is created

```

[acadgild@localhost ~]$ hadoop fs -ls /
18/03/24 12:38:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 11 items
drwxr-xr-x - acadgild supergroup 0 2018-03-04 22:22 /Prachi
drwxr-xr-x - acadgild supergroup 0 2018-03-24 09:38 /hadoopdata
drwxr-xr-x - acadgild supergroup 0 2018-03-04 19:15 /input
drwxr-xr-x - acadgild supergroup 0 2018-03-05 09:06 /output
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:22 /output1
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:32 /output2
drwxr-xr-x - acadgild supergroup 0 2018-03-24 12:37 /output3
drwxr-xr-x - acadgild supergroup 0 2018-03-17 09:59 /sqoopout
drwxr-xr-x - acadgild supergroup 0 2018-02-02 12:49 /sqoopout111
drwxrwx--- - acadgild supergroup 0 2018-03-05 08:40 /tmp
drwxr-xr-x - acadgild supergroup 0 2018-03-18 07:10 /user
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ █

```

```

[acadgild@localhost ~]$ hadoop fs -ls /output3
18/03/24 12:39:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-03-24 12:37 /output3/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 30 2018-03-24 12:37 /output3/part-r-00000
[acadgild@localhost ~]$ hadoop fs -cat /output3/part-r-00000
18/03/24 12:39:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count 17
length 78
square 424
[acadgild@localhost ~]$ █

```

